

Generating Uniform Spanning Trees

Haonan Duan, Yangjun Ruan

In this lecture, we introduce two randomized algorithms for generating spanning trees uniformly. Given a connected, undirected and unweighted graph $G(V, E)$ with n vertices and m edges, the goal is to generate a spanning tree T of G chosen **uniformly** at random among all spanning trees of G . Recall that a spanning tree T of an undirected graph G is a subgraph that is a tree which includes all vertices of G .

1 The first algorithm: random walk

A random walk on G is a Markov chain over all vertices with the transition matrix:

$$P(w|v) = \begin{cases} \frac{1}{r_v} & \text{if } (v, w) \text{ is an edge.} \\ 0 & \text{o.w.} \end{cases}$$

where r_v is the degree of the node v .

ALGORITHM I [ALD90] [BRO89]

Simulate a simple random walk on the graph G starting at an arbitrary vertex s until every vertex is visited. For each vertex $i \in V - s$ collect the edge $\{j, i\}$ that corresponds to the first entrance to vertex i . Let T be this collection of edges. Output T .

Our first algorithm is that a random walk can simulate a uniform spanning tree by selecting the first edge that enters the node. The pseudocode is shown above. The running time of Algorithm I is $O(mn)$. It's clear that T is a spanning tree because it contains $|V| - 1$ edges and no loops. We shall prove that it is indeed uniformly distributed.

Before discussing the algorithm, it will be useful to introduce *arborescences* defined as below:

Definition 1.1. (Arborescence) For a given $s \in G$, an arborescence T rooted at s is a directed spanning tree of G where all vertices in $G \setminus \{s\}$ have *exactly one* incoming arc.

It is easy to check that there is a one-to-one correspondence between spanning trees of G and arborescences rooted at s : given any spanning tree, there is a unique way of directing its edges to make it an arborescence rooted at s ; conversely, given any arborescence rooted at s , one can obtain a spanning tree by simply disregarding the direction of the edges. Therefore, if we get a procedure that generates arborescences, it also gives a procedure that generates spanning trees.

Indeed, random walk simulates a chain of arborescences along the way.

Definition 1.2. (Arborescence chain). Let $(X_j; -\infty < j < \infty)$ be the Markov chain of a random walk on G . We can define a Markov chain of arborescences $(S_j; -\infty < j < \infty)$. Considering

(X_j, X_{j+1}, \dots) , we choose X_j to be the root of S_j and select the first edge entering other vertices to be in the tree. In other words, the tree S_j is rooted at X_j with edges $(X_{T_v^j-1}, X_{T_v^j}); v \neq X_j$, where $T_v^j = \min\{k \leq j; X_k = v\}$. An example is given in Figure 1.

0	1	2	3	4	5	6	7	...
1	2	1	3	4	1	3	2	...




Figure 1: An example of an arborescence chain. The table above is a random walk over graph G . Then by Definition 1.2, $S_0 = \{[1, 2], [1, 3], [1, 4]\}$ and $S_3 = \{[3, 2], [3, 4], [3, 1]\}$

Proposition 1.3. $(S_j; -\infty < j < \infty)$ is an irreducible and reversible Markov chain.

Proof. Irreducibility and reversibility are directly from the assumption that G is connected and unweighted. \square

Theorem 1.4. Let $N(G)$ be the number of spanning trees t of G . Then $P(T = t) = 1/N(G)$ for each t produced by Algorithm 1.

We will first prove a simpler case when the graph is r -regular.

Proof. Let G be an r -regular graph. Let $(X_j; -\infty < j < \infty)$ and $(S_j; -\infty < j < \infty)$ be the Markov chain of the nodes and arborescences introduced previously.

Consider the transition matrix Q for S_j in reversed time:

$$Q(t, t') = P(S_{-1} = t' | S_0 = t)$$

Then

- Given t , there are exactly r trees t' such that $Q(t, t') = \frac{1}{r}$ and $Q(t, t') = 0$ for other trees.
- Given t' , there are exactly r trees t such that $Q(t, t') = \frac{1}{r}$ and $Q(t, t') = 0$ for other trees.

This is because X_{-1} has $\frac{1}{r}$ probability to be each of the r neighbours of v . Each possibility leads to S_{-1} being some tree t' , and these also are the only possibilities. A similar argument can be made going from S_{-1} to S_0 . Therefore, Q is doubly-stochastic, which further implies that the stationary distribution of S_j is uniform. \square

To extend the above proof to non-regular graphs, we first make the observation that each tree t has $r(v)$ neighbours, where v is the root of the tree. This implies that we can swap r in the two bullet points above with the degree of the roots. It then follows that the stationary distribution of t is proportional to the degree of its root. Thus, all directed spanning trees rooted at the same node are uniform. Since our trees are undirected and unrooted, all T in Algorithm 1 are uniform.

2 Faster Generation of Spanning Trees

In this section, we will introduce a faster algorithm for generating *approximately uniform* random spanning trees. In particular, we focus on the generation of δ -random spanning trees:

Definition 2.1 (δ -random spanning trees). A randomized algorithm A that generates δ -random spanning trees outputs a random spanning tree T with probability $p(T)$ that is δ -far from uniform, i.e.,

$$\frac{1 - \delta}{|\mathcal{T}(G)|} \leq p(T) \leq \frac{1 + \delta}{|\mathcal{T}(G)|}$$

where $\mathcal{T}(G)$ is the set of spanning trees of G .

With the same reasoning as before, if we get a procedure that generates δ -random arborescences, it also gives a procedure that generates δ -random spanning trees.

[KM09] introduces a faster generation algorithm which generates δ -random spanning trees in expected time of $\tilde{O}(m\sqrt{n}\log(1/\delta))$. For brevity of illustration, we will focus on a simplified algorithm that gives an expected running time of $\tilde{O}(m^2/\sqrt{n}\log(1/\delta))$, which shares the same spirit.

The key idea of the algorithm is the observation that the random walk algorithm may spend a long time walking in regions that have already been covered. Indeed, the random walk algorithm has a running time of $O(mn)$, while only a tiny fraction $O(n)$ is used for recovering an arborescence. Therefore, the algorithm seeks to obtain a shortcut that cuts out the random walk corresponding to visiting already explored parts of G . The essential steps are to first decompose the graph into small subgraphs that can be quickly covered, and then shortcut the walk inside each subgraph *if it is already covered*.

2.1 Decompose the graph

First, we define the (ϕ, γ) -Decomposition that will permit the implementation of the fast generation algorithm. Let (D_1, \dots, D_k, S, C) denote a partition of G , where D_i are disjoint subgraphs, $S = V(G) \setminus \bigcup_i V(D_i)$ is the set of remaining vertices, and $C = E(G) \setminus \bigcup_i E(D_i)$ is the set of edges not entirely contained inside one of D_i . For a given D_i , let $C(D_i)$ be the subset of C incident to D_i and $U(D_i)$ be the set of vertices of D_i incident to an edge from C .

Definition 2.2 ((ϕ, γ) -decomposition). (D_1, \dots, D_k, S, C) is a (ϕ, γ) -decomposition if:

1. $|C| \leq \phi|E(G)|$
2. $\forall i$, the diameter $\gamma(D_i) \leq \gamma$
3. $\forall i$, $|C(D_i)| \leq |E(D_i)|$

Note that the first condition ensures that the subgraphs D_1, \dots, D_k contain most (all but a ϕ fraction of) edges in G . Intuitively, the second condition ensures that each subgraph could be covered relatively quickly, using the fact that the cover time of an unweighted graph G' with diameter $\gamma(G')$ is at most $O(|E(G')|\gamma(G'))$ [Ale+79].

For reasons that will become clear later, we consider a specific decomposition that can be quickly computed, shown by the following lemma:

Lemma 2.3 (Obtaining good $((\phi, \gamma)$ -decompositions). *For G and any $\phi = o(1)$, a $(\phi, \tilde{O}(1/\phi))$ -decomposition of G can be computed in time $\tilde{O}(m)$.*

We omit the proof for brevity and refer interested readers to Lemma 13 and its proof in [KM09].

2.2 Decompose the walk

Then, we consider the random walk $X = (X_i)$ over the decomposed graph started from a vertex chosen by the stationary distribution G . Let τ be the cover time of G , i.e., the first time that the walk visits all the vertices of G . The result in [Ale+79] yields that $E[\tau] = O(mn)$ which is the expected time of traditional random walk.

We decompose the walk over decomposed G . Let Z and Z_i be the random variables corresponding to the number of times that X traverses edges from C and inside D_i respectively. By definition, we have that $\tau = \sum_i Z_i + Z$. The following lemma establishes the expectation of Z :

Fact 2.4. *The expected traversed time in C is $E(Z) = O(\phi mn)$.*

Proof. Since the random walk starts from the stationary distribution of G , the expected number of traversals of edges from C is just proportional to its size. Therefore, the assumption $|C| \leq \phi|E(G)|$ implies the above fact. \square

As discussed before, we will be interested in the cover time for each subgraph. In particular, let Z_i^* be the random variable corresponding to the number of traversals inside D_i until X explores the whole subgraph D_i , we have that:

Lemma 2.5 (Expected cover time for subgraphs). $E[Z_i^*] = \tilde{O}(|E(D_i)|\gamma(D_i))$.

Proof. Let us fix $D = D_i$. For a vertex $v \in V(D)$, let $d_G(v)$ be the degree of v in G and $d_D(v)$ be the degree of v in D . For $u, v \in U(D)$, let $p_{u,v}^D$ be the probability that a random walk in G that starts at u will reach v through a path that does not pass through any edge inside D . Consider a (weighted) graph D' , which we obtain from D by adding, for each $u, v \in U(D)$, an edge (u, v) with weight $d_G(u)p_{u,v}^D$. We note that if we take our walk X and filter out of the vertices that are not from D , then the resulting "filtered" walk Y_D will just be a natural random walk in D' . As a result, it is easy to see that, in this case, $E[Z_i^*]$ can be upper-bounded by the expected time needed by a random walk in D' , started at arbitrary vertex, to visit all of the vertices in D' and then reach some vertex in $U(D)$. Therefore, to bound $E[Z_i^*]$, it suffices to bound the cover time of D' . The cover time of undirected graph G' is at most $2 \log |V(G')| H_{\max}(G')$, where $H_{\max}(G')$ is the maximal hitting time. Following the results of [Ale+79], we have $H_{\max}(G') \leq |E(G')|\gamma(G')$. \square




Figure 2: (a) The shortcut of trajectory inside D_1 is implemented by $P_v(e)$ that characterizes the probability of X leaving D_1 through edge e after entering through vertex v ; (b) The illustration for computing $P_v(e)$ by adding a dummy node u^* which connects to all other leaving edges.

2.3 Shortcut the walk

Recall that the key idea of the algorithm is to shortcut the trajectory inside D_i after D_i has already been covered. More specifically, consider the walk after D_i is covered, this means that we already know for all edges $v \in V(D_i)$ which arc e_v should be added to the arborescence. Any further walk inside D_i provides no new information for constructing the arborescence, and could be shortcut.

The main tool for implementing the shortcut is $P_v(e)$ that characterizes the probability of X leaving D_i through edge e after entering through vertex v , see Fig. 2a for an illustration. If we know $P_v(e)$ for all $v \in V(D_i)$ and all $e \in C(D_i)$, then we could essentially immediately choose the leaving edge e upon entering v without computing the explicit trajectory in D_i .

To formalize the intuition, let $\bar{X} = (\bar{X}_1, \dots, \bar{X}_m)$ be a decomposition of the walk $X = (X_1, \dots, X_\tau)$ into contiguous blocks \bar{X}_j that are contained in D_{i_j} for $i_j \in \{0, \dots, k\}$, where $D_0 = S$. We construct the shortcutted walk $\tilde{X} = (\tilde{X}_i)$ by processing \bar{X} block by block: if D_{i_j} has not been covered or $i_j = 0$, we copy \bar{X}_j to \tilde{X} , otherwise we copy only the first and last entries of the block.

With $P_e(v)$, the following lemma establishes that \tilde{X} can be simulated efficiently:

Lemma 2.6 (Simulating \tilde{X} with $P_v(e)$). *Knowing $P_v(e)$ for all $e \in C(D_i), v \in V(D_i)$ and i , we can preprocess these values in $\tilde{O}(\phi mn)$ time and it allows simulation of l steps of \tilde{X} in time $\tilde{O}(l)$.*

Proof. Simulating \tilde{X} before D_i is covered is straightforward. The only thing to show is that the shortcutting of blocks can be simulated efficiently. To show this, we consider some i and $v \in U(D_i)$, and construct an array $A_v(m)$ where $A_v(m) = \sum_{1 \leq j \leq m} P_v(e_j)$ for $m \in \{1, \dots, |C(D_i)|\}$ in $\tilde{O}(|C(D_i)|)$. So we can construct for all $v \in U(D_i)$ in $\tilde{O}(|C(D_i)||V(D_i)|)$. Summing over all D_i gives the above bound on preprocessing time. Furthermore, this array enables choosing e with $P_v(e)$ with binary search in polylogarithmic time. \square

From this lemma, we can show that we can find a random arborescence of G efficiently:

Lemma 2.7 (Finding a random arborescence). *Given a (ϕ, γ) -decomposition of G and $P_v(e)$, we can find a random arborescence of G in expected time $\tilde{O}(m(\gamma + \phi n))$.*

Proof. We only need to compute the expected length of the shortcutted walk \tilde{X} , which is upper-bounded by the following quantity:

$$\sum_i \underbrace{E[Z_i^*]}_{\text{cover time of } D_i} + \underbrace{E[Z]}_{\text{traversals in } C} + \underbrace{2E[Z]}_{\text{shortcutted traversals}}$$

The first and the second terms are obvious, and the last term is due to the fact that the two vertices that remain in \tilde{X} after shortcutting some block from X can be amortized into the number of traversals by X of some edges in C . By the second assumption in Definition 2.2, Fact 2.4, Lemma 2.5, we get that $\sum_i E[Z_i^*] + 3E[Z] = \tilde{O}(\sum_i |E(D_i)|\gamma + \phi mn) = \tilde{O}(m(\gamma + \phi n))$. By Lemma 2.6, the shortcutted work can be simulated in expected time $\tilde{O}(m(\gamma + \phi n))$. \square

Then the remaining question is whether we can compute $P_e(v)$ efficiently, established by the following lemma:

Lemma 2.8 (Computing $P_v(e)$). *Given a (ϕ, γ) -decomposition of G , we can compute multiplicative $(1 + \varepsilon)$ -approximations of $P_v(e)$ in time $\tilde{O}(\phi m^2 \log(1/\varepsilon))$*

Proof. Let us fix some $D = D_i$ and an edge $e = (u, u') \in C(D)$ with $u \in U(D)$. Consider a graph D' obtained by adding a dummy node u^* to D and then connecting all other leaving edges to u^* , i.e., for each $(w, w') \in C(D) \setminus \{e\}$, we add an edge (w, u^*) (note that w' can be equal to u'), see Fig. 2b for an illustration. Note that $P_v(e)$ is the probability that the walk started at v will *hit* u' *before* u^* , which can be quickly computed using *electrical flows*.

In particular (see, e.g., [Lov93]), we can treat D' as an electrical circuit where we impose voltage of 1 at u' and 0 at u^* , then the voltage achieved at v is exactly equal to $P_v(e)$. We can compute a $(1 + \varepsilon)$ -approximation in $\tilde{O}(|E(D')| \log 1/\varepsilon)$ using the linear solver in [ST14]. Computing for each $e \in C(D)$ and store the probabilities for all vertices v , the total running time can be bounded by $\tilde{O}(|C| \sum_i |E(D_i)| \log 1/\varepsilon) = \tilde{O}(\phi m^2 \log 1/\varepsilon)$, where we use the fact that $|E(D')| = |E(D)| + |C(D)| \leq 2|E(D)|$ by assumption 3 in Definition 2.2. \square

Note that the above lemma only gives an approximation of $P_v(e)$. We need to show that it is sufficient for controlling the overall error and maintaining a good running time:

Lemma 2.9 (An approximate $P_v(e)$ is sufficient). *Given a (ϕ, γ) -decomposition of G and multiplicative $(1 + \varepsilon)$ -approximation of $P_v(e)$, we can generate a δ -random arborescence of G in expected time $\tilde{O}(m^2(\gamma + \phi n))$ as long as $\varepsilon \leq \delta/mn$.*

We omit the proof for brevity and refer interested readers to Lemma 10 and its proof in [KM09].

Putting all together, we have the theorem establishing the running time of the algorithm:

Theorem 2.10 (Total complexity). *For any $\delta > 0$, we can generate a δ -random spanning tree of G in expected time of $\tilde{O}(m/\sqrt{n} \log(1/\delta))$.*

Proof. Let $\phi = 1/n^{1/2}$, $\varepsilon = \delta/mn$, the total expected time can be decomposed as:

1. Get a $(1/n^{1/2}, \tilde{O}(n^{1/2}))$ -decomposition: $\tilde{O}(m)$ (Lemma 2.3).
2. Compute estimate of $P_v(e)$: $\tilde{O}(m^2/\sqrt{n} \log(1/\delta))$ (Lemma 2.8).
3. Generate a δ -arborescence: $\tilde{O}(m\sqrt{n})$ (Lemma 2.7, Lemma 2.9).

Summing all together, we have the total running time of $\tilde{O}(m^2/\sqrt{n} \log(1/\delta))$. \square

We refer interested readers to Sec. 4 in [KM09] for an improved algorithm with $\tilde{O}(m\sqrt{n} \log(1/\delta))$, which is obtained by a stronger decomposition and a slightly different random walk.

References

- [Ald90] David J Aldous. “The random walk construction of uniform spanning trees and uniform labelled trees”. In: *SIAM Journal on Discrete Mathematics* 3.4 (1990), pp. 450–465 (cit. on p. 1).
- [Ale+79] Romas Aleliunas, Richard M Karp, Richard J Lipton, László Lovász, and Charles Rackoff. “Random walks, universal traversal sequences, and the complexity of maze problems”. In: *20th Annual Symposium on Foundations of Computer Science (sfcs 1979)*. IEEE Computer Society. 1979, pp. 218–223 (cit. on p. 4).
- [Bro89] Andrei Z Broder. “Generating random spanning trees”. In: *FOCS*. Vol. 89. 1989, pp. 442–447 (cit. on p. 1).
- [KM09] Jonathan A Kelner and Aleksander Madry. “Faster generation of random spanning trees”. In: *2009 50th Annual IEEE Symposium on Foundations of Computer Science*. IEEE. 2009, pp. 13–21 (cit. on pp. 3, 4, 6, 7).
- [Lov93] László Lovász. “Random walks on graphs”. In: *Combinatorics, Paul erdos is eighty* 2.1-46 (1993), p. 4 (cit. on p. 6).
- [ST14] Daniel A Spielman and Shang-Hua Teng. “Nearly linear time algorithms for preconditioning and solving symmetric, diagonally dominant linear systems”. In: *SIAM Journal on Matrix Analysis and Applications* 35.3 (2014), pp. 835–885 (cit. on p. 6).